Ashwin Chembu

(341) 356-1722 | ashchembu@gmail.com | San Francisco, California | https://ashchembu.com

PROFILE

Data engineer building Snowflake SQL pipelines, healthcare entity-matching workflows, and traceable quality controls. Connects layered ETL/MDM architecture with API-assisted reconciliation and stakeholder-ready evidence.

TECHNICAL SKILLS

Languages: SQL, Python, TypeScript, R

Data Platforms: Snowflake, MongoDB, AWS S3, AWS Glue, Pandas, NumPy

Data Engineering: Layered ETL/MDM, entity resolution, identifier matching, delta logic

Quality / Delivery: Reconciliation, data profiling, source-to-target mappings, QC reports, stewardship review

EXPERIENCE

Business Technology Solutions Associate | ZS Associates **South San Francisco, CA | Jun 2024-Present**

- Built **Snowflake SQL** pipelines over **327K+** records for identifier matching, address comparison, and quality reporting.
- Designed matching across **67K+** records using healthcare identifiers and fuzzy-search signals for stewardship review.
- Defined layered **ETL/MDM** flows from landing to publication with delta logic, history tracking, and data-quality checks.
- Translated stakeholder needs into testable rules, API integrations, source-to-target mappings, and delivery documents.

Founder and Lead Developer | Tallyrus **Jun 2023-Sep 2025**

- Built Flask services with **MongoDB** persistence and S3 storage to manage documents and support essay-feedback flows.
- Built **Python**/LangChain essay-feedback workflows with Pinecone and OpenAI for retrieval-augmented educator review.
- Integrated React/TypeScript and Express flows with Google OAuth, **MongoDB,** sessions, and protected product routes.

Data Analyst / Web Development Lead | Tech4Good at UCSC **Mar 2023-Jun 2025**

- Analyzed **43,340** interaction rows across **2,805** sessions with **Python**/Pandas and built Excel/XLSX exports for research.
- Separated chat, text, and poll state in three **MongoDB** databases, using change streams to synchronize edits and votes.
- Implemented 20+ attributed Node/Express routes with **MongoDB** persistence and logging for chat, text entry, and polls.

PROJECTS

Forge | Personal Fitness Intelligence **Personal project**

- Integrated Hevy/Oura APIs and Apple Health webhooks with FastAPI/**MongoDB** to unify training, nutrition, and recovery.
- Implemented configurable Hevy/Oura synchronization, on-demand refresh, and CSV ingestion to keep fitness records current.
- Compared performance with a 12-week hypertrophy plan to score workout adherence, nutrition compliance, and recovery.

Azure RAG | Document Search Proof of Concept **Public proof of concept**

- Built **Python** document search with Azure OpenAI embeddings, LangChain, and FAISS to retrieve relevant source passages.
- Added PDF/TXT/Markdown ingestion, chunking, vector indexing, and source citations to ground interactive document Q&A.
- Documented Azure embedding/chat-model setup and a 30+ document corpus for repeatable retrieval and grounded Q&A.

Lume | AI Meeting Intelligence **Dec 2025-Jan 2026**

- Built React and Node/Express flows with **MongoDB**/S3 to capture, store, search, and query meetings and project files.
- Built reusable Project Brain knowledge bases retaining transcript/file terminology and relationships for grounded Q&A.
- Designed RAG with vector retrieval, intent routing, evidence extraction, and generator/critic passes for cited answers.

EDUCATION

University of California, Davis | **B.S. Computer Science** | **GPA: 3.7** Graduated Jun 2025

CERTIFICATIONS

AWS Certified Cloud Practitioner | Google Data Analytics Certificate

LINKS

Portfolio: https://ashchembu.com | LinkedIn: linkedin.com/in/ashwin-chembu-a5292b205 | GitHub: github.com/ashwinchembu